

Complete Guide - Graph RAG Agent Workflow with n8n

Ahmed Aziz Ammar

Contents

1	Introduction	3
2	Prerequisites	3
3	Docker Desktop Installation	3
4	Installing n8n in Self-Hosted Mode with Docker	3
5	Accessing the n8n Interface	4
6	Configuring Neo4j in the Cloud	5
7	Installing Neo4j in n8n	6
8	Overview of the Graph RAG Workflow	7
8.1	General Architecture	8
9	Importing the First Workflow into n8n	9
10	Workflow No. 1: Graph Database Workflow	9
10.1	Node 1.1: Execute Workflow	10
10.2	Node 1.2: Read/Write Files from Disk	10
10.3	Node 1.3: Extract from File	11
10.4	Node 1.4: Edit Fields1	11
10.5	Node 1.5: Execute Command	12
10.6	Code Node 1.6: Extract All Information (!!Important)	13
10.7	Code Node 1.7: Bad Format Cleanup	13
10.8	Node 1.8: Edit Fields	14
10.9	Important: Creating the Neo4j Credential	15
10.10	Filling in the Connection Fields	16
10.11	Code Node 1.9 : Create Batch + Cypher Query	17
10.12	Node 2.0: Create Database (Neo4j)	18
10.13	Code Node 2.1 : Return Items Separated	19
10.14	Relationship Nodes: Creating Relationships → AI Agent → Bidirectional Pairs → ExecuteQuery graphDb	20

11 Click the Execute Workflow Button	27
12 Import the Second Workflow into n8n	27
13 Workflow No. 2: Graph RAG Agent	28
13.1 Architecture	28
13.2 Node 2.1: When Chat Message Received	28
13.3 Node 2.2: AI Agent	28
13.4 Node 2.3: OpenAI Chat Model	30
13.5 Node 2.4: Postgres Chat Memory	30
13.6 Node 2.5: ExecuteQuery graphDb in Neo4j	31
14 Test the Graph RAG	32
15 Conclusion	33

1 Introduction

This guide details the steps to install and run a Graph RAG workflow locally with a cloud-based graph database, using **n8n**. n8n is a no-code/low-code automation tool for orchestrating tasks such as document analysis, API calls, file management, etc.

2 Prerequisites

Before getting started, make sure you have the following elements on your machine:

- GIT installed
- 2.7 GB space for Docker Desktop
- 4.5 GB space for self-hosted-ai-starter-kit container
- Terminal or Command Prompt

3 Docker Desktop Installation

1. Download Docker Desktop from the official link: <https://www.docker.com/products/docker-desktop>
2. Install Docker Desktop following the instructions for Windows.
3. Verify that Docker is operational with the command:

```
docker --version
```

4 Installing n8n in Self-Hosted Mode with Docker

1. Clone the official self-hosted-ai-starter-kit repository: `git clone https://github.com/n8n-io/self-hosted-ai-starter-kit.git` Check if there's a space between "https" and ":" — if so, concatenate them.
2. Access the cloned directory: `cd self-hosted-ai-starter-kit`
3. Copy the example .env file and rename it: `copy .env.example .env`
4. Launch the installation with Docker Compose:

```
docker-compose --profile cpu up
```

Check if there's a space between "-" and "profile" — if so, concatenate them.

5. Check the location of the `self-hosted-ai-starter-kit` folder on your machine. We need the `.env` file for PostgreSQL configuration and the `shared` folder for resource location.

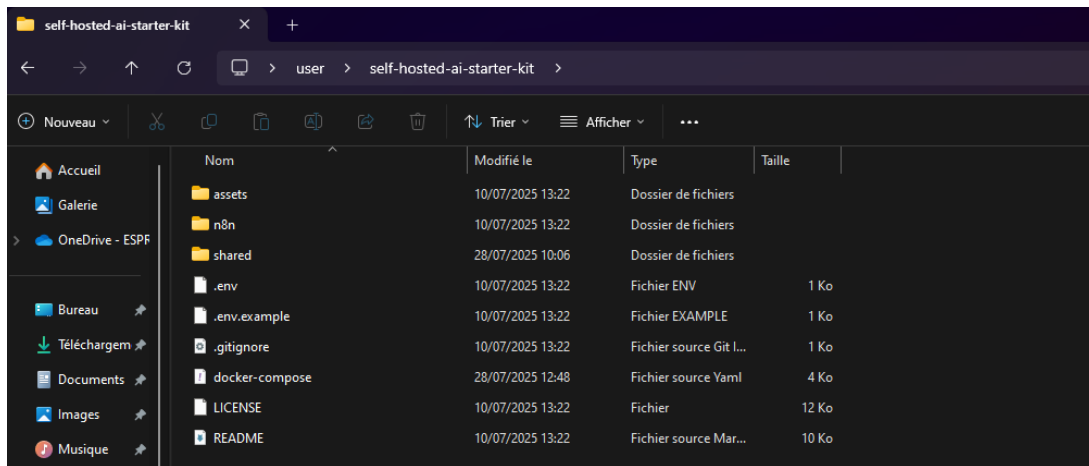


Figure 1: Structure of the self-hosted-ai-starter-kit folder

5 Accessing the n8n Interface

1. After installation, open Docker Desktop and verify that the container `self-hosted-ai-starter-kit` is running.
2. Click the generated link: `5678:5678` to access the web interface, or open: <http://localhost:5678>

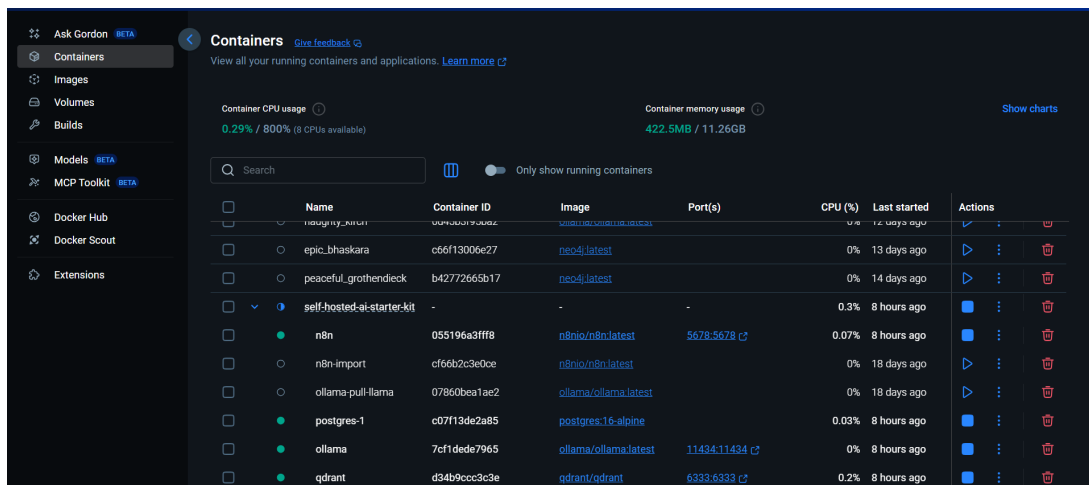


Figure 2: self-hosted-ai-starter-kit container in Docker Desktop

3. Complete the registration form in n8n (First name, Last name, Email, etc.).

References:

- Official n8n guide: <https://github.com/n8n-io/self-hosted-ai-starter-kit>
- Download Docker Desktop: <https://www.docker.com/products/docker-desktop>

6 Configuring Neo4j in the Cloud

To set up Neo4j in cloud mode, we use the **Neo4j AuraDB** platform, which hosts graph databases and provides secure API access.

- Visit the official website: <https://console.neo4j.io/>.
- Sign in using your Gmail or Outlook account.
- Once logged in, you will be redirected to an interface prompting you to download a file containing your credentials:
 - **NEO4J_USERNAME**
 - **NEO4J_PASSWORD**
 - **NEO4J_URI**
- Download and store this file securely. Wait a few minutes for the free instance to be created (named “free instance” in Neo4j Aura).

```
# Wait 60 seconds before connecting using these details, or login to https://console.neo4j.io to validate the Aura Instance is available
NEO4J_URI=neo4j+s://bf4fb4b6.databases.neo4j.io
NEO4J_USERNAME=neo4j
NEO4J_PASSWORD=veatHt8y9mh1LuU9VtNkLmXwoM3dm-pFddyYS8w
NEO4J_DATABASE=neo4j
AURA_INSTANCEID=bf4fb4b6
AURA_INSTANCENAME=trial instance
```

Figure 3: Neo4j AuraDB credentials file

- You will then be redirected to the Neo4j Aura management interface.

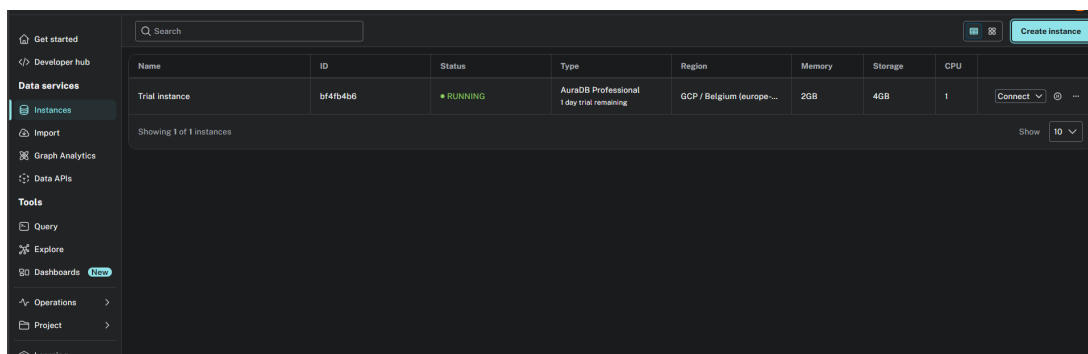


Figure 4: Neo4j Aura interface

- Click on “Connect” and select “Query”. You will be redirected to the Cypher Query interface.




Figure 5: Neo4j connection options

- You can execute Cypher commands to visualize nodes and their relationships.
- **Example query:**

```
MATCH (n:YourGraphDatabaseName) RETURN n LIMIT 50;
```

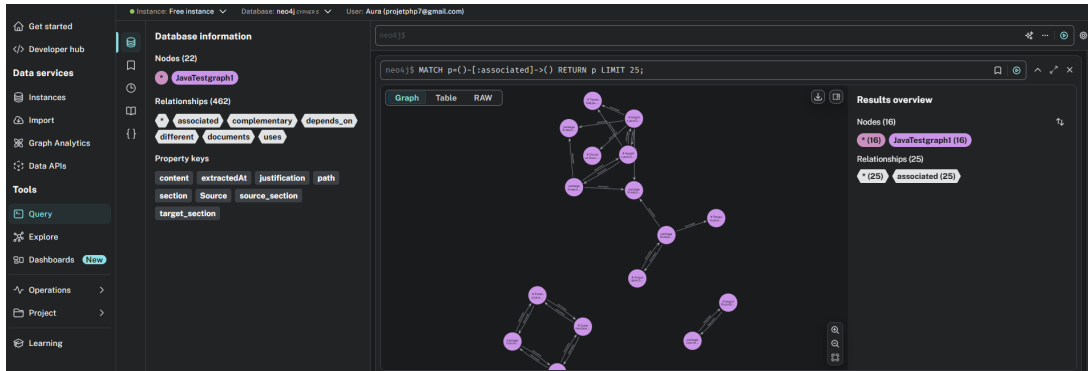


Figure 6: Visualizing relationships between nodes in Neo4j

7 Installing Neo4j in n8n

To integrate Neo4j with n8n, follow these steps:

1. Go to **Settings** in n8n.
2. Select the **Community Nodes** tab.
3. Click the **Install** button.
4. In the search bar, type: `n8n-nodes-neo4j`.
5. Select the package found, then click **Install**.

Illustrative Screenshots

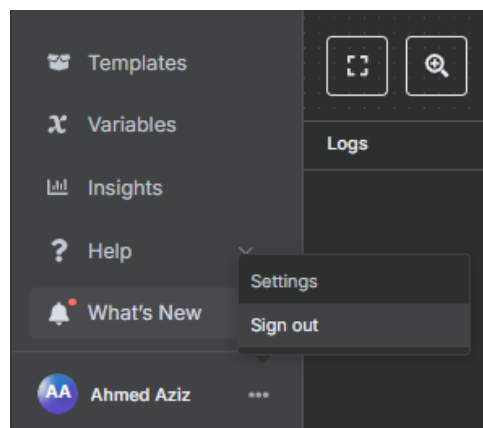


Figure 7: Settings menu in n8n

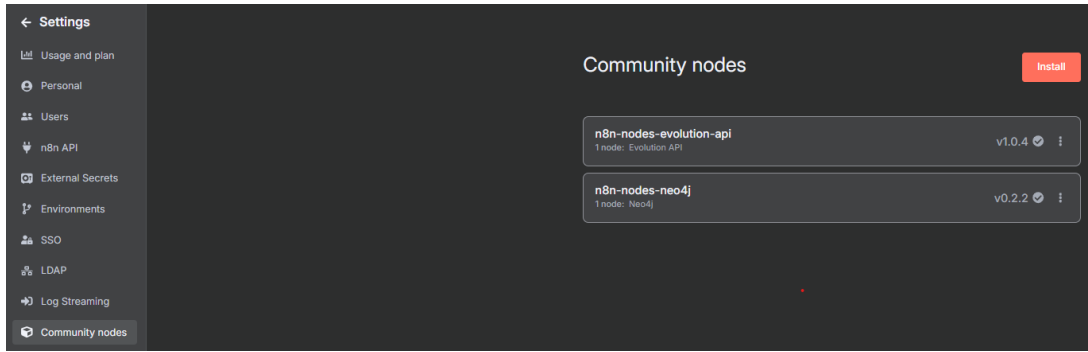


Figure 8: Community Nodes tab for extension installation

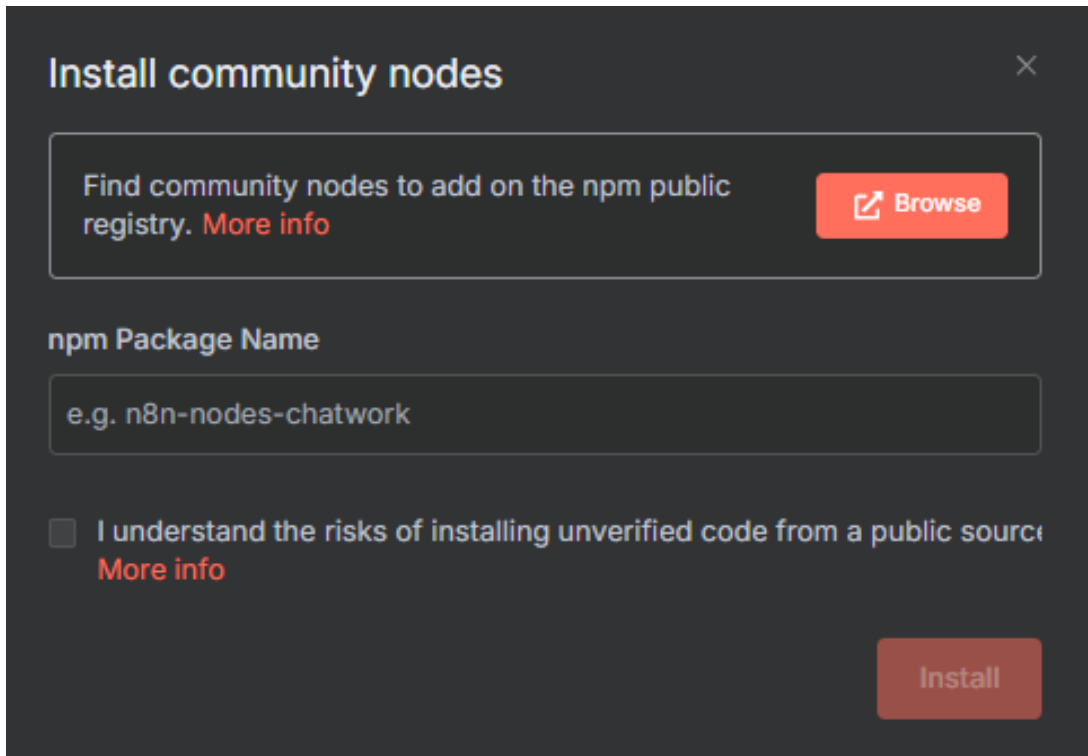


Figure 9: Installing the Neo4j node (n8n-nodes-neo4j)

8 Overview of the Graph RAG Workflow

These two workflows serve the following main purposes:

- Automate the ingestion of documents (Java, Markdown) from local directories.
- Index and store these documents as nodes in a cloud-hosted Neo4j graph database.
- Automatically generate semantic relationships between documents (e.g., *documents*, *complementary*, *different*).
- Enable advanced queries and synthetic responses based on the graph (**Graph RAG** approach).

- Provide justified answers with metadata from source documents (title, section, extraction date, content snippet).

Figure 10: Graph RAG Agent Workflow

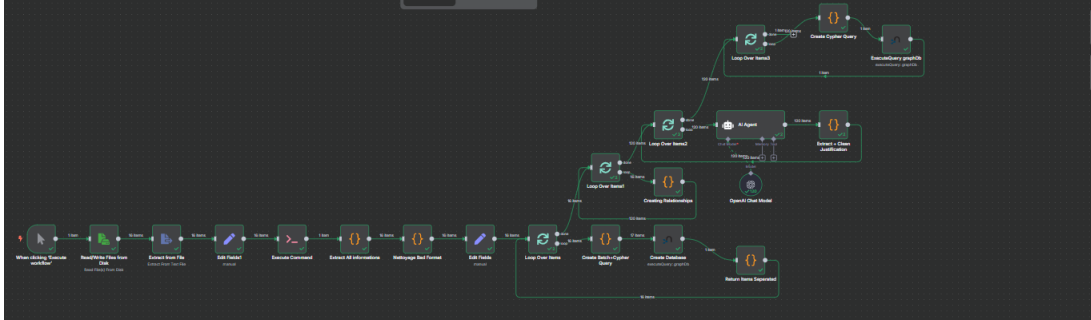


Figure 11: Ingestion and setup workflow for Neo4j graph database

8.1 General Architecture

The architecture is based on two main interconnected workflows built around the Neo4j graph database:

1. Graph Construction Workflow This first workflow aims to ingest and structure knowledge from source files (Java, Markdown). It includes:

- **Reading and extracting documents:** retrieving file content and metadata.
- **Cleaning and normalization:** preparing data to avoid insertion errors.
- **Creating nodes and relationships in Neo4j:** inserting files as enriched nodes and automatically generating semantic relationships between them.

2. Graph RAG Agent Workflow This second workflow enables conversational interaction with the graph:

- **Receiving user queries** through a chat interface.
- **Retrieving relevant information** by querying Neo4j with a RAG agent.
- **Multi-document reasoning** to provide a synthesized response based on files and their relationships, without hallucination.
- **Conversational memory** to ensure continuity in dialogue.

Knowledge Base (Neo4j) Neo4j is the core of the architecture, organizing documents and their links as a graph, enabling advanced searches and contextual reasoning.

9 Importing the First Workflow into n8n

1. Launch n8n
2. Go to <http://localhost:5678>
3. Navigate to the Workflows menu > Create Workflow




Figure 12: n8n home interface

4. Click on > Import from File




Figure 13: Import from a file

5. Select the file Graph Database Resources:JAVA+MD.json
6. Click Save

10 Workflow No. 1: Graph Database Workflow

Node Architecture

Manual Trigger → Read Files → Extract Content → [Edit Fields1] → Execute Command → Extract All Info → Clean Bad Format → [Edit Fields] → [Loop Over Items(100 items)] → Create Batch+Cypher → Create Database → Return Items → Creating Relationships → [Loop Over Items1(1000 items)] → [Loop Over Items2(5000 items)] → AI Agent + OpenAI Model → Extract+Clean → [Loop Over Items3 (5K)] → Create Cypher Query → Create Relationships NEO4J

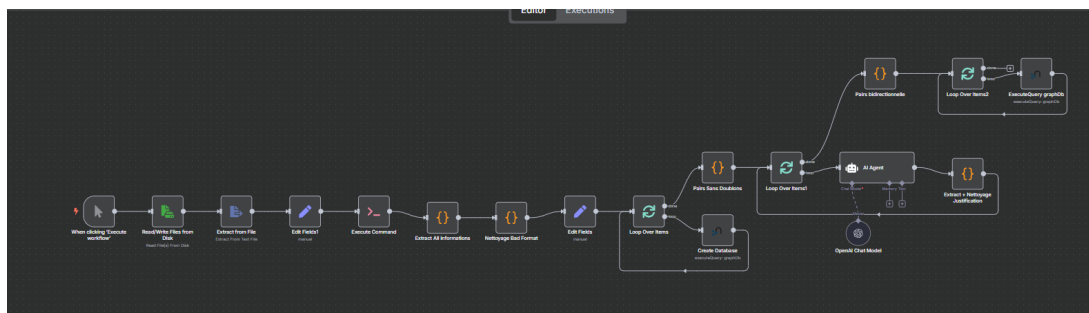


Figure 14: Graph creation workflow in Neo4j

Node Configuration

10.1 Node 1.1: Execute Workflow

Role: Manually triggers the file indexing and relationship generation process.

Configuration:

- **Type:** Manual Trigger
- **Usage:** Click to start the workflow that reads the files, cleans the content, and inserts the data into Neo4j.

10.2 Node 1.2: Read/Write Files from Disk

Role: Recursively reads files from the local filesystem.

Configuration:

- **Operation:** Read Files
- **File(s) Selector:** To read all ‘.java’ and ‘.md’ files, including from subfolders, make sure n8n is installed via Docker. Add your resources to the **shared** directory inside **self-hosted-ai-starter-kit**:
`/data/shared/YourPATH/**/*.java,md.JAVA,md,markdown`

Explanation of the glob pattern:

- ****:** Recursively traverses subdirectories
- *****: Any file
- **{JAVA,md,markdown}**: Targeted file extensions

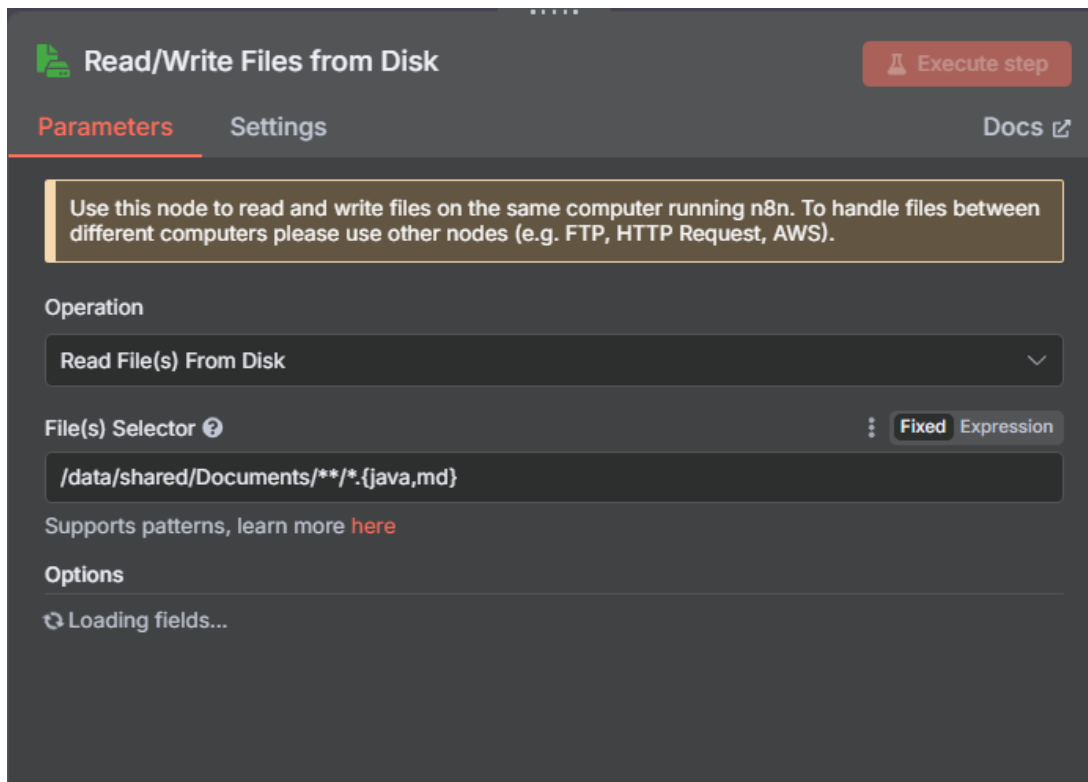


Figure 15: File reading configuration

10.3 Node 1.3: Extract from File

Role: Extracts raw content from files (TXT, MD, Markdown, Java).

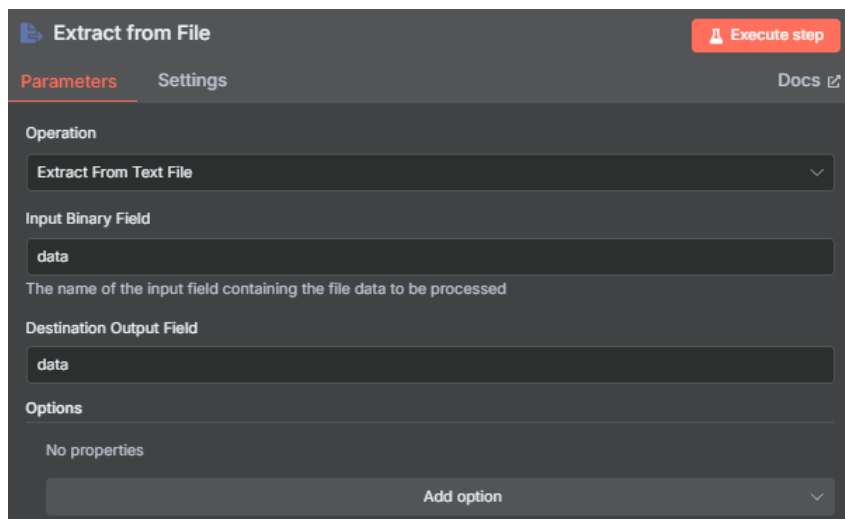


Figure 16: File extraction configuration

10.4 Node 1.4: Edit Fields1

Role: Normalizes the extracted fields to create properties such as `file_name` and `content`.

Edit Fields1 Execute step

Parameters Settings Docs

Mode
Manual Mapping

Fields to Set

file_name	A String	= {{ \$('Read/Write Files from Disk').item.json.fileName }}
content	A String	= {{ \$json.data }}

Drag input fields here or Add Field

Include Other Input Fields ☐

Options
No properties
Add option

Figure 17: Field editing configuration

10.5 Node 1.5: Execute Command

Role: Lists all file paths using the shell command:

```
find /data/shared/YourPATH/*
```

Execute Command Execute step

Parameters Settings Docs

Execute Once ☒ Fixed Expression

Command Fixed Expression

```
find /data/shared/resource/*
```

Figure 18: Shell command execution configuration

10.6 Code Node 1.6: Extract All Information (!!Important)

Role: Associates each detected file path with the extracted metadata and content in preparation for insertion into Neo4j.

!!WARNING: It is crucial to ****modify line 32 of the code**** to specify the ****root path of your resources folder****. Without this change, no data will be correctly linked in the graph database.

Example:

```
const relativePath = fullPath.replace('/data/shared/resource/', '');
```

Tip: This path must point to your `/data/shared/` directory because it is the mounted volume used by n8n via Docker.

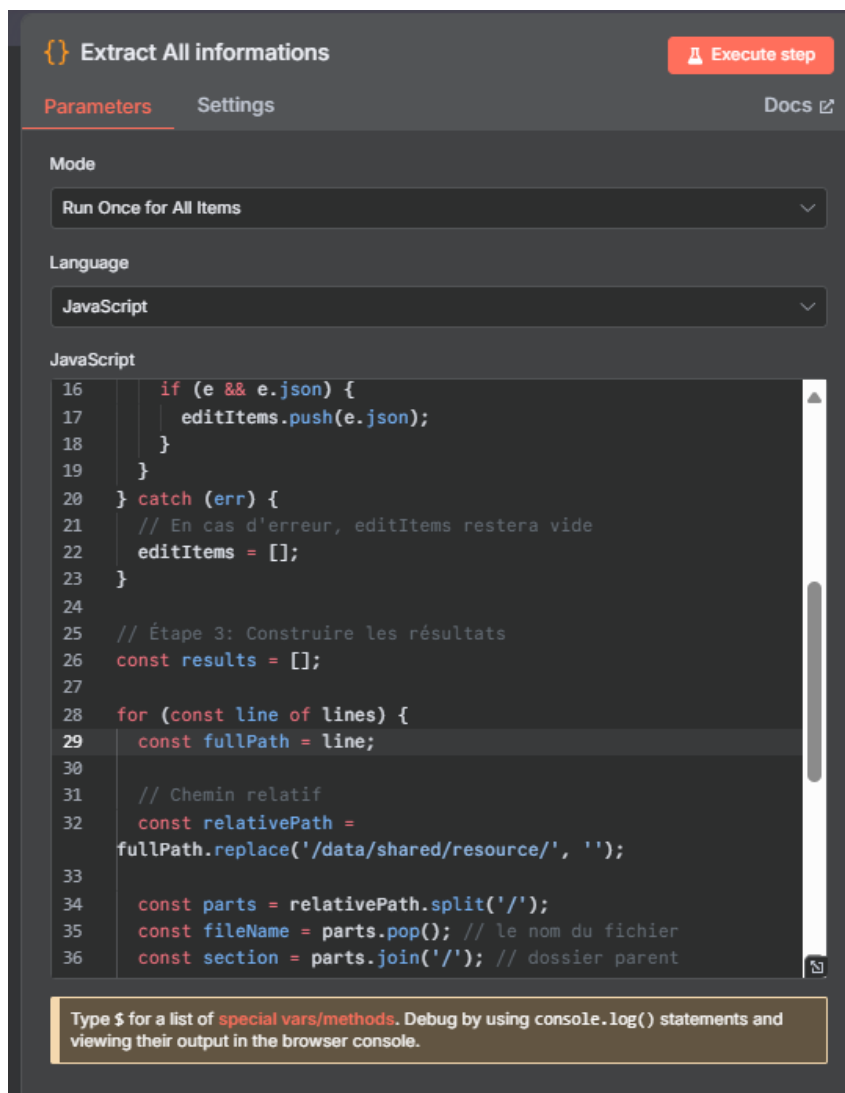


Figure 19: Code for extracting information

10.7 Code Node 1.7: Bad Format Cleanup

Role: Cleans the extracted content by removing unwanted characters (e.g., “”, line breaks) and escaping special characters before inserting into Neo4j.

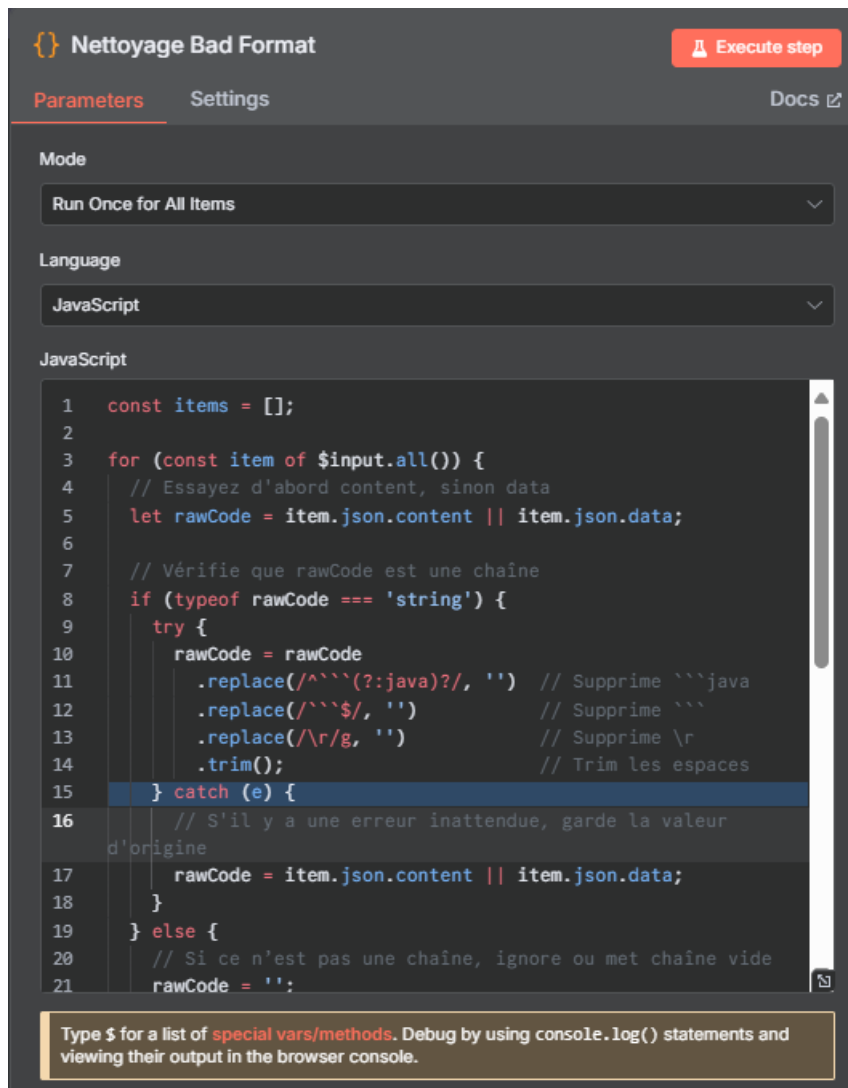



Figure 20: Format cleanup code

10.8 Node 1.8: Edit Fields

Role: Prepares final properties for Neo4j (file name, section, path, cleaned content).

 Edit Fields

Execute step

Parameters

Settings

Docs

Mode

Manual Mapping

Fields to Set

file_name

A String

=

{{ \$json.file }}

section

A String

=

{{ \$json.section }}

PATH

A String

=

{{ \$json.path }}

content

A String

=

{{ \$json.cleanedCode }}

extractedAT

A String

=

{{ \$json.extractedAt }}

Drag Input fields here or Add Field

10.9 Important: Creating the Neo4j Credential

In the first Neo4j node, click the first field + **Create new Credential**.

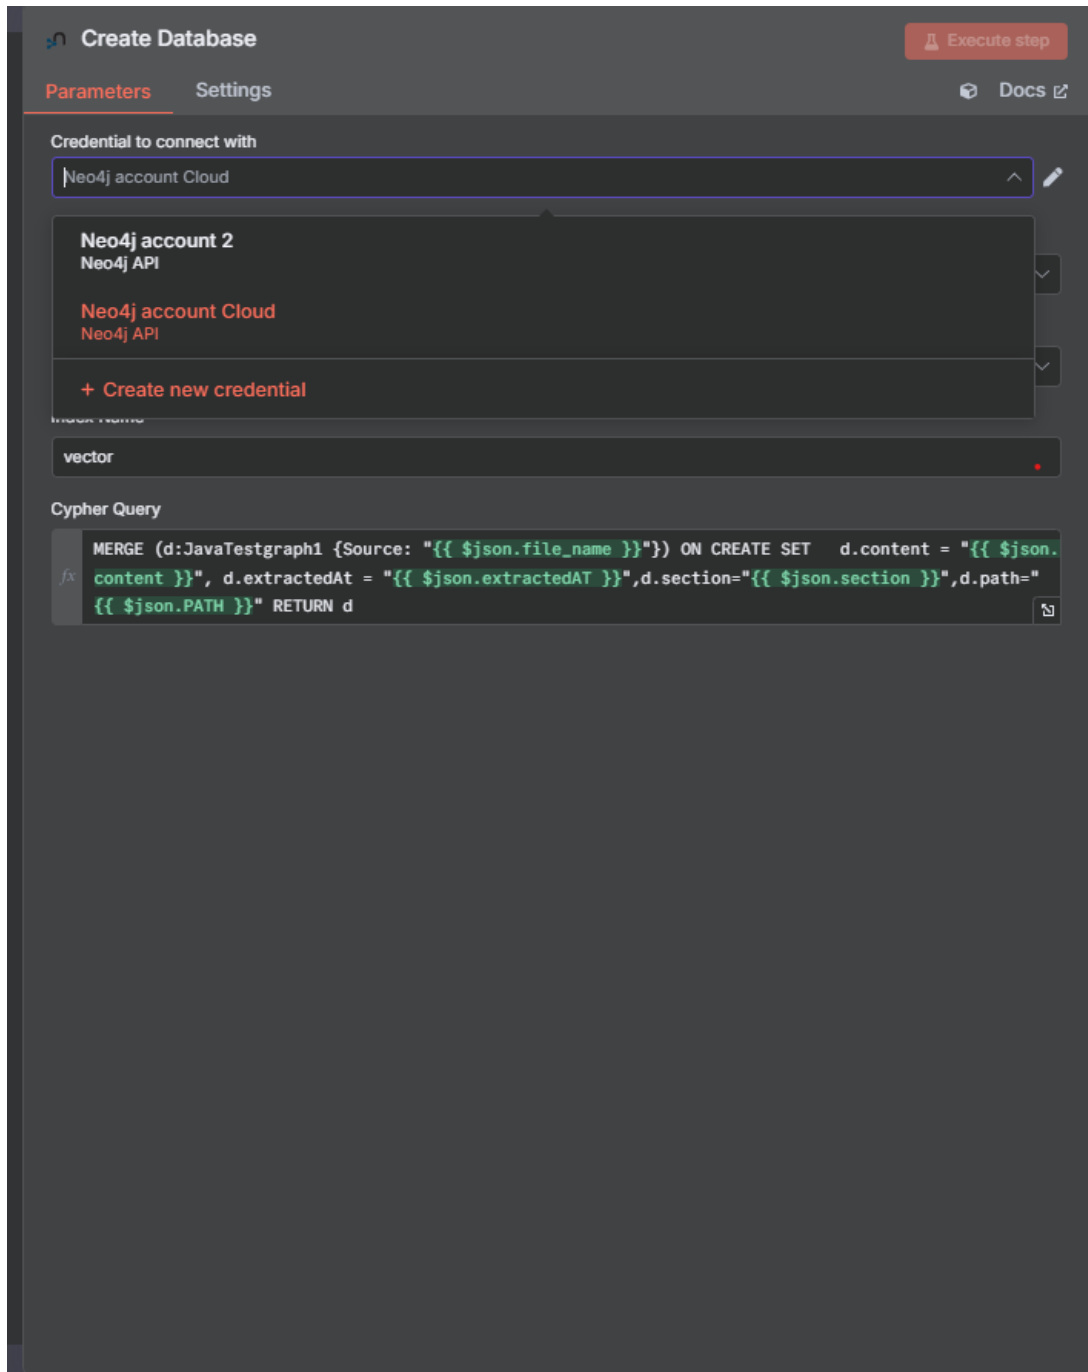


Figure 22: Creating a Neo4j credential in n8n

10.10 Filling in the Connection Fields

Fill in the following fields using the values from the credentials file you downloaded when creating the Neo4j Aura instance:

- **Host:** NEO4J_URI (e.g., neo4j+s://xxxx.databases.neo4j.io)
- **Username:** NEO4J_USERNAME (typically neo4j)
- **Password:** NEO4J_PASSWORD (generated during instance creation)

- **Database:** NEO4J_DATABASE (often neo4j)

The screenshot shows the 'Neo4j account Cloud' configuration window. On the left, there are tabs for 'Connection', 'Sharing', and 'Details'. The 'Connection' tab is active. A yellow banner at the top says 'Need help filling out these fields? Open docs'. The form contains the following fields:

- Connection URI ***: neo4j+s://9c42a465.databases.neo4j.io
- Username ***: neo4j
- Password ***: (masked with dots)
- Database ***: neo4j

At the bottom, a small information icon is followed by the text: 'Enterprise plan users can pull in credentials from external vaults. More info'.

Figure 23: Neo4j connection configuration in n8n

Important: Once created, the credential will automatically be used in other Neo4j nodes, and the same logic applies to other n8n nodes.

10.11 Code Node 1.9 : Create Batch + Cypher Query

Role: Processes batches of files (up to 100 items per batch) and generates optimized Cypher queries for bulk insertion into Neo4j database. This node also preserves all original file items for downstream processing.

Key Functions:

- **escapeCypher():** Sanitizes string values to prevent Cypher injection
- **Batch processing:** Handles up to 100 files simultaneously
- **Error handling:** Skips invalid items with warnings
- **Content truncation:** Limits content to 8000 characters per field

Important Note: On line 36 of the code, users must modify the graph database name (`TestGraph1`) according to their specific database configuration requirements.

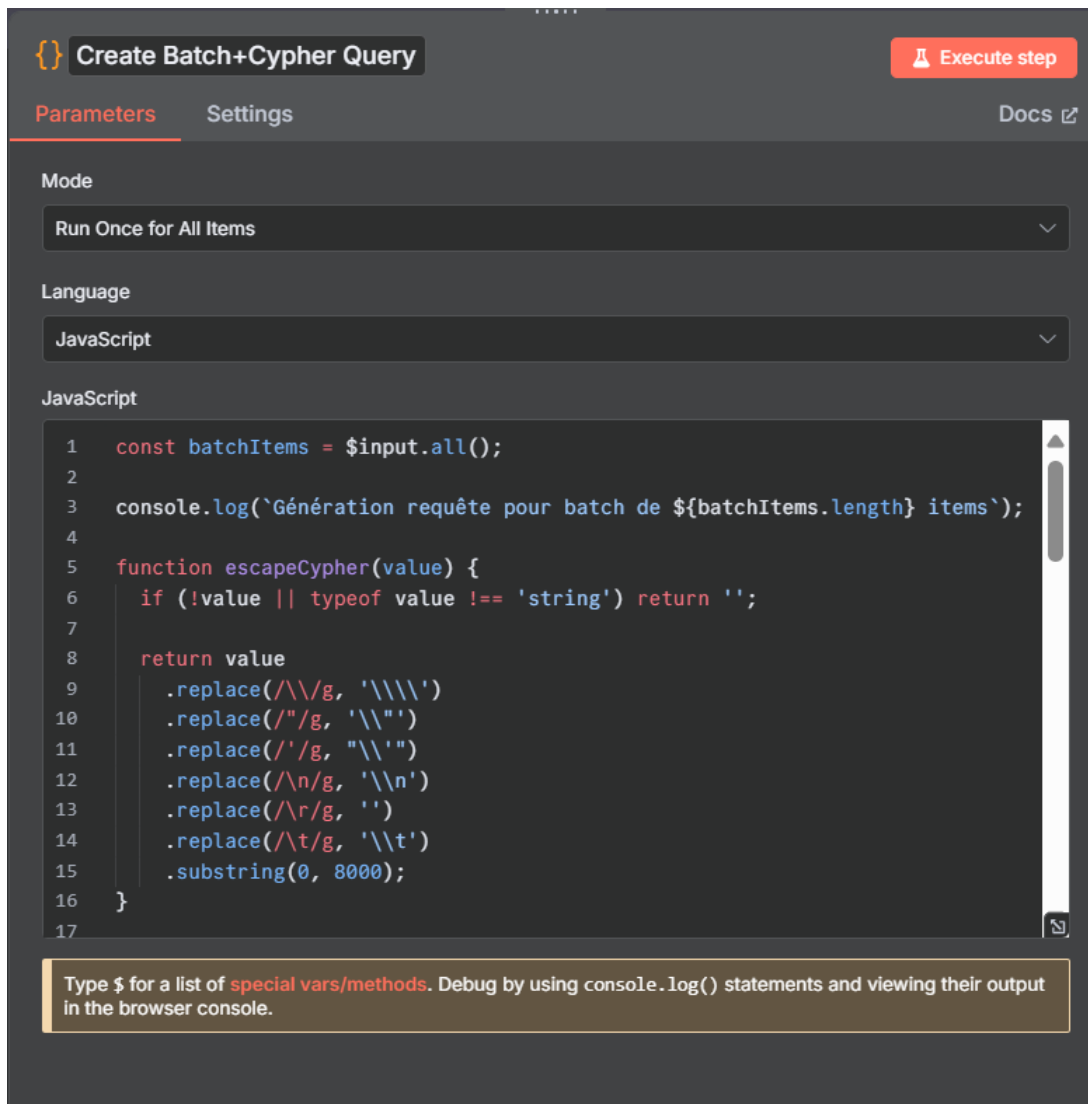


Figure 24: Code node configuration for batch Cypher query generation in N8N

10.12 Node 2.0: Create Database (Neo4j)

Role: Inserts nodes (files) into the Neo4j database with the following properties:

- **Source:** File name
- **content:** Cleaned content
- **section:** Parent folder
- **path:** Full file path

Query Configuration:

- **Resource:** Graph Database
- **Operation:** Execute Query
- **Cypher Query :**Output of Previous Node Code

```
{{ $('Create Batch+Cypher Query').item.json.cypherQuery }}
```

The screenshot shows the 'Create Database' configuration window in Neo4j. It has a dark theme and a top bar with the title 'Create Database' and an 'Execute step' button. Below the title bar are two tabs: 'Parameters' (selected) and 'Settings'. The 'Parameters' tab contains several sections: 'Credential to connect with' with a dropdown menu showing 'Neo4j account Cloud'; 'Resource' with a dropdown menu showing 'Graph Database'; 'Operation' with a dropdown menu showing 'Execute Query'; 'Index Name' with a text input field containing 'vector'; and 'Cypher Query' with a text input field containing the JSON path query `{{ $('Create Batch+Cypher Query').item.json.cypherQuery }}`. There is also a 'Docs' link in the top right corner.

Figure 25: Execute Query node for creating nodes

10.13 Code Node 2.1 : Return Items Separated

Role: Retrieves and separates the original file items that were preserved during the batch processing workflow. This node ensures that all processed files are available individually for subsequent relationship creation between file pairs, eliminating duplicates and maintaining data integrity for the next phase of the knowledge graph construction.

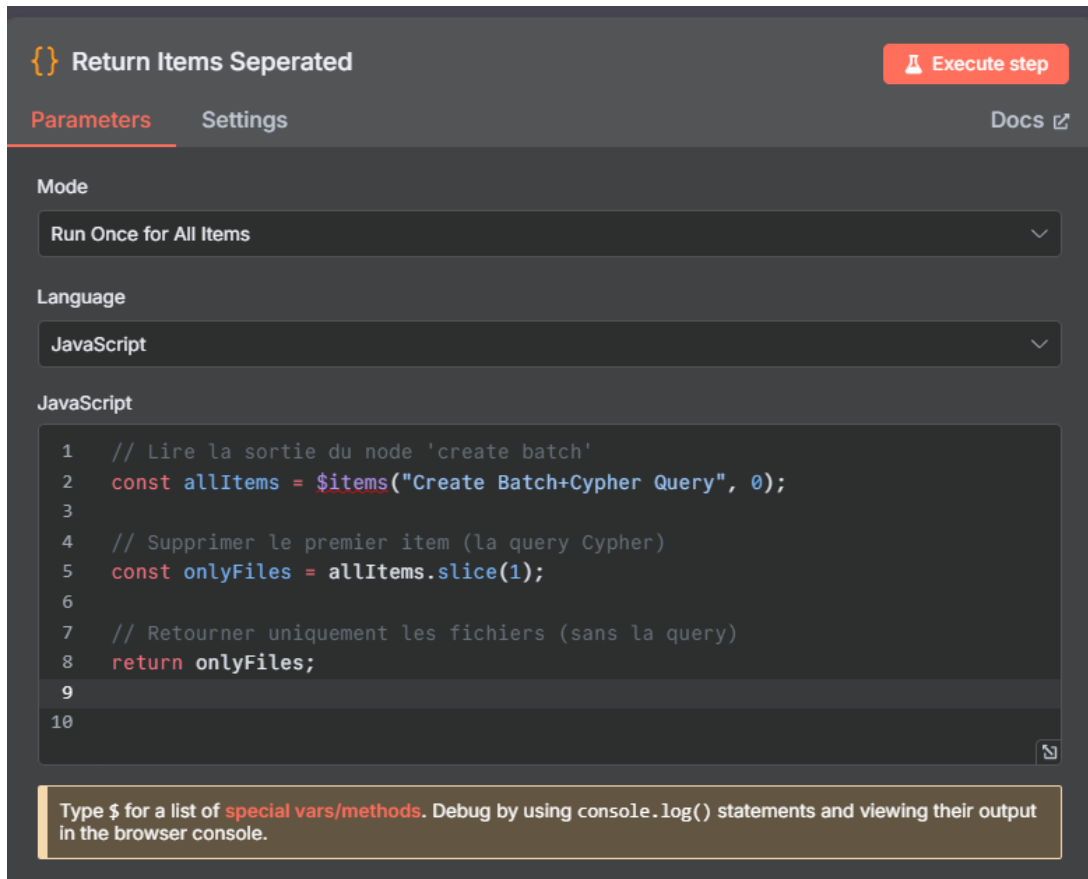


Figure 26: Code node configuration for separating and returning original file items

10.14 Relationship Nodes: Creating Relationships → AI Agent → Bidirectional Pairs → ExecuteQuery graphDb

Overall Role: Build semantic relationships between files based on their content and sections, and insert them into Neo4j.

Processing Chain:

1. **Creating Relationships(Code Node)** Generates all file pairs (A, B) without repetition (**no** $A \rightarrow A$ and **no inverse duplicates** $B \rightarrow A$). Each pair includes:
 - **source:** {file, path, section, content}
 - **target:** {file, path, section, content}

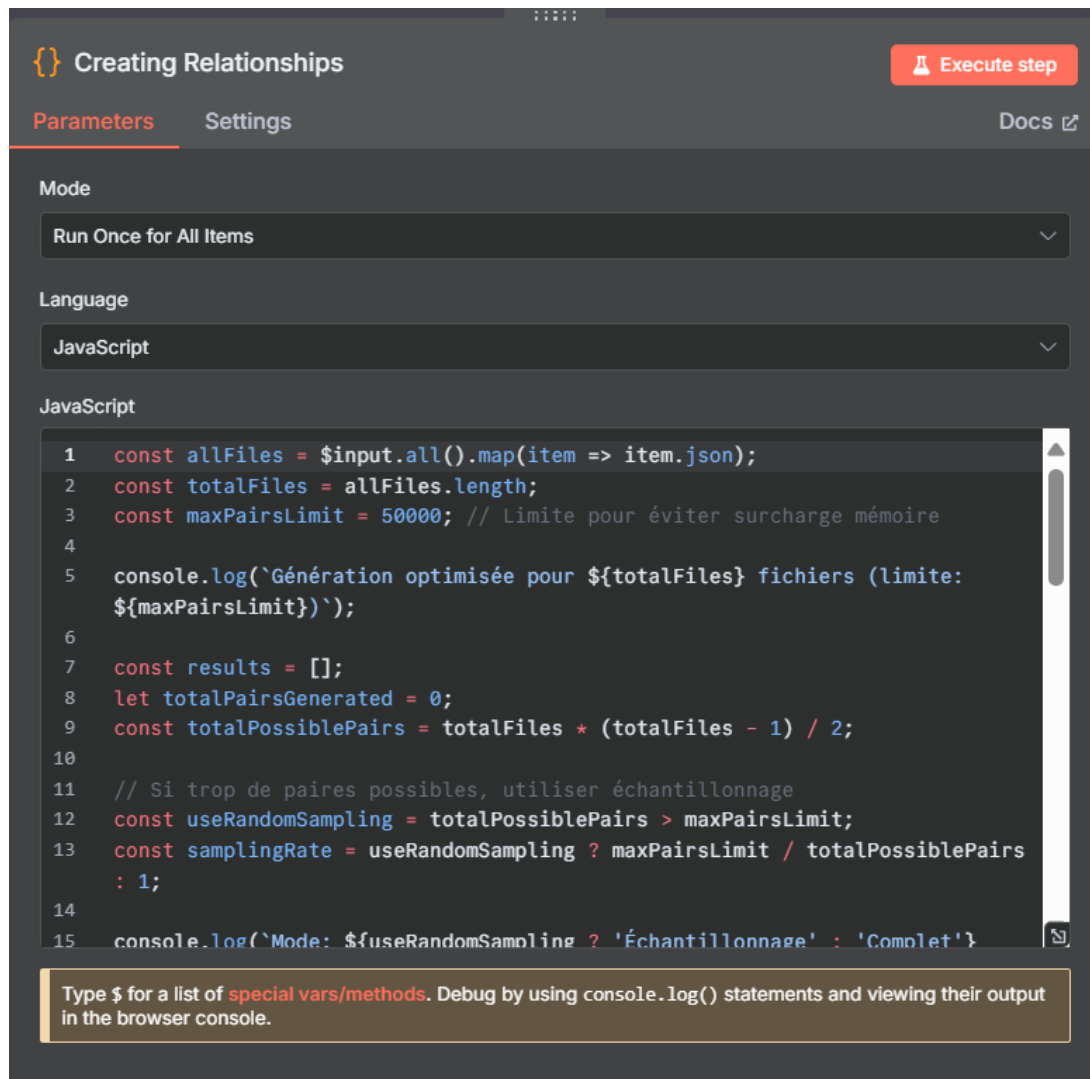


Figure 27: Code Node generating unique pairs

2. **AI Agent (LangChain)** Analyzes each (*source*, *target*) pair to determine:

- type: Relationship type (complementary, uses, documents, associated, etc.)
- justification: Concise explanation of the relationship
- source_section, target_section
- source, target (file names)

The image shows the 'AI Agent' configuration interface. At the top, there's a header with 'AI Agent' and an 'Execute step' button. Below the header, there are tabs for 'Parameters' and 'Settings'. The 'Parameters' tab is active. Under 'Source for Prompt (User Message)', there's a dropdown menu set to 'Define below'. The 'Prompt (User Message)' section contains a text area with the following content: 'Analyze the following two files and define a relevant and general relationship between them. Source: File: {{ \$json.source.file }} Path: {{ \$json.source.path }}'. Below the prompt, there are two toggle switches: 'Require Specific Output Format' and 'Enable Fallback Model', both currently turned off. The 'Options' section has a 'System Message' text area with the content: 'You are an expert in source code analysis, technical document classification, and knowledge graph modeling. Your task is to identify relevant and general relationships between two files by analyzing their content, section, and technical nature. Important rules: You are an expert in source code analysis, technical document classification, and knowledge graph modeling. Your task i...'. At the bottom of the options section, there's an 'Add Option' button.

Figure 28: AI agent configuration for relationship analysis

3. OpenAI GPT 4.1 mini Chat Model (Configuration)

- Go to the OpenAI Embeddings Node and click + **Create new Credentials**

The image shows the 'OpenAI Chat Model' configuration interface. At the top, there's a header with the OpenAI logo and 'OpenAI Chat Model'. Below the header, there are tabs for 'Parameters' and 'Settings'. The 'Parameters' tab is active. Under 'Credential to connect with', there's a dropdown menu with the text 'Select Credential' and a red warning icon. Below the dropdown, there's a button labeled '+ Create new credential'. The 'Options' section has a 'Sampling Temperature' input field with the value '0,0'. At the bottom of the options section, there's an 'Add Option' button.

Figure 29: Creating a connection with OpenAI

- Fill in the connection fields using your OpenAI API Key: <https://platform.openai.com/api-keys>

OpenAI account

Connection

Sharing

Details

Connection tested successfully Retry

Need help filling out these fields? [Open docs](#)

API Key *

Organization ID (optional)

Only required if you belong to multiple organisations

Base URL

<https://api.openai.com/v1>

Enterprise plan users can pull in credentials from external vaults. [More info](#)

Figure 30: OpenAI Connection Fields

4. **Code Node Extract + Clean Justification (Batch Processing)** This node processes all items from a batch (typically 100 items) in a single execution and has three main roles:
 - (a) **Batch Processing:** Uses `$input.all()` to retrieve all items from the current batch instead of processing one item at a time, enabling efficient bulk processing within the Loop Over Items2 iteration.
 - (b) **JSON Extraction & Parsing:** For each item in the batch:
 - Extracts raw AI model output from `item.json.output`
 - Removes Markdown code blocks (“`json`” and “`”`)
 - Parses the cleaned JSON string into a structured object
 - Handles parsing errors gracefully with error objects
 - (c) **Structured Output Generation:** Creates clean objects containing:
 - `source` and `target`: file names from parsed JSON
 - `source_section` and `target_section`: directory paths
 - `relationshipType`: extracted from `parsed.type`
 - `justification`: cleaned and Cypher-safe text
 - `extractedAt`: ISO timestamp for traceability
 - `batchIndex`: item position for debugging

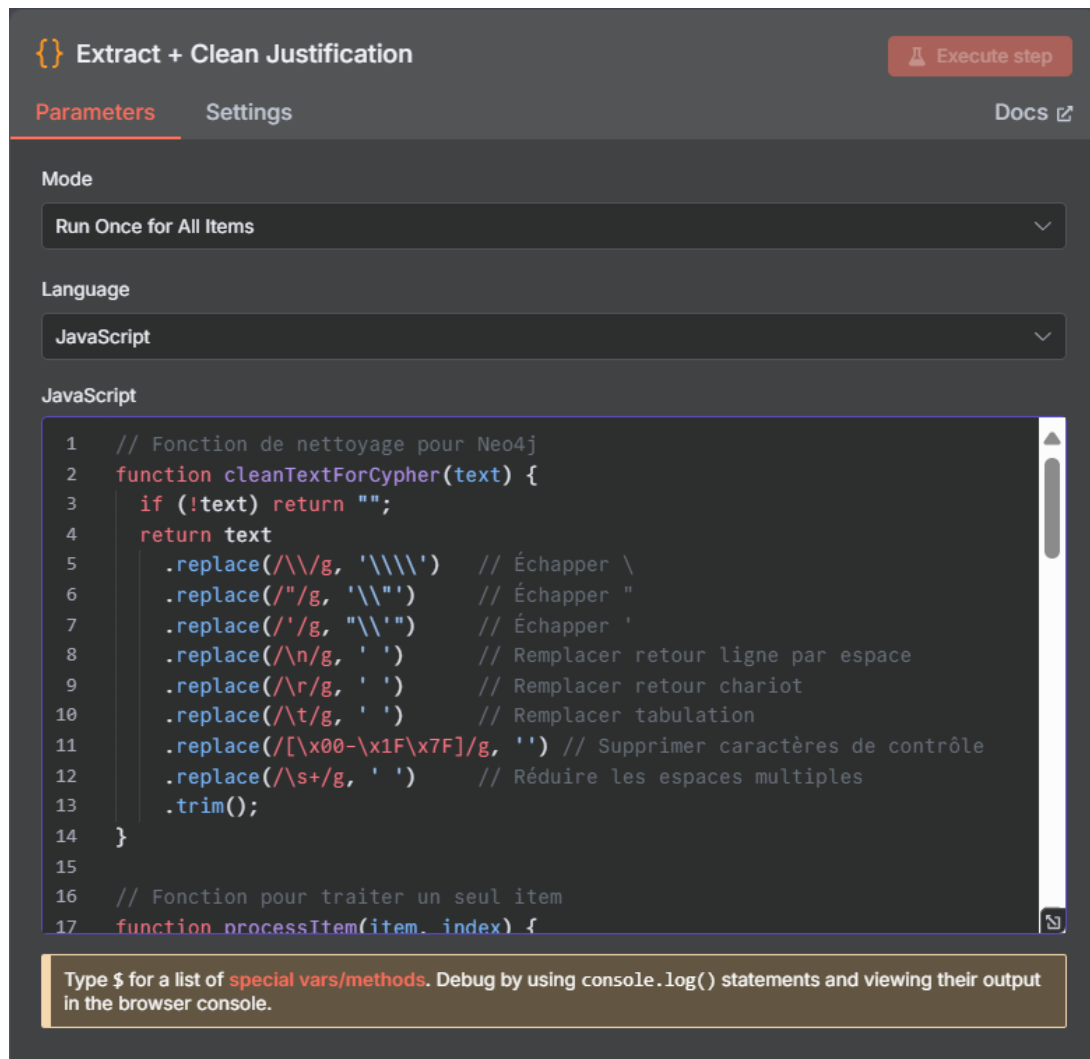


Figure 31: Code Node for extraction and cleaning of justifications

- (d) **Code Node Create Cypher Query (Batch MERGE Generator) Role** This node processes batches of relationships (typically 5000 items from Loop Over Items3) and generates optimized Cypher queries for bulk Neo4j insertion:
- i. **Batch Processing:** Retrieves all relationships from the current batch using `$input.all()` to process multiple relations in a single query execution.
 - ii. **Advanced Cypher Escaping:** Implements `escapeCypher()` function with enhanced security:
 - Escapes special characters: `\`, `"`, `'`
 - Converts newlines to `\n` and tabs to `\t`
 - Truncates values at 8000 characters to prevent query overflow
 - Handles null/undefined values gracefully
 - iii. **Dynamic MERGE Statement Generation:** For each relationship in the batch:
 - Validates required fields (source, target, relationshipType)
 - Creates unique node variables (`src${i}`, `tgt${i}`) to avoid conflicts
 - Generates MERGE statements for both source and target nodes

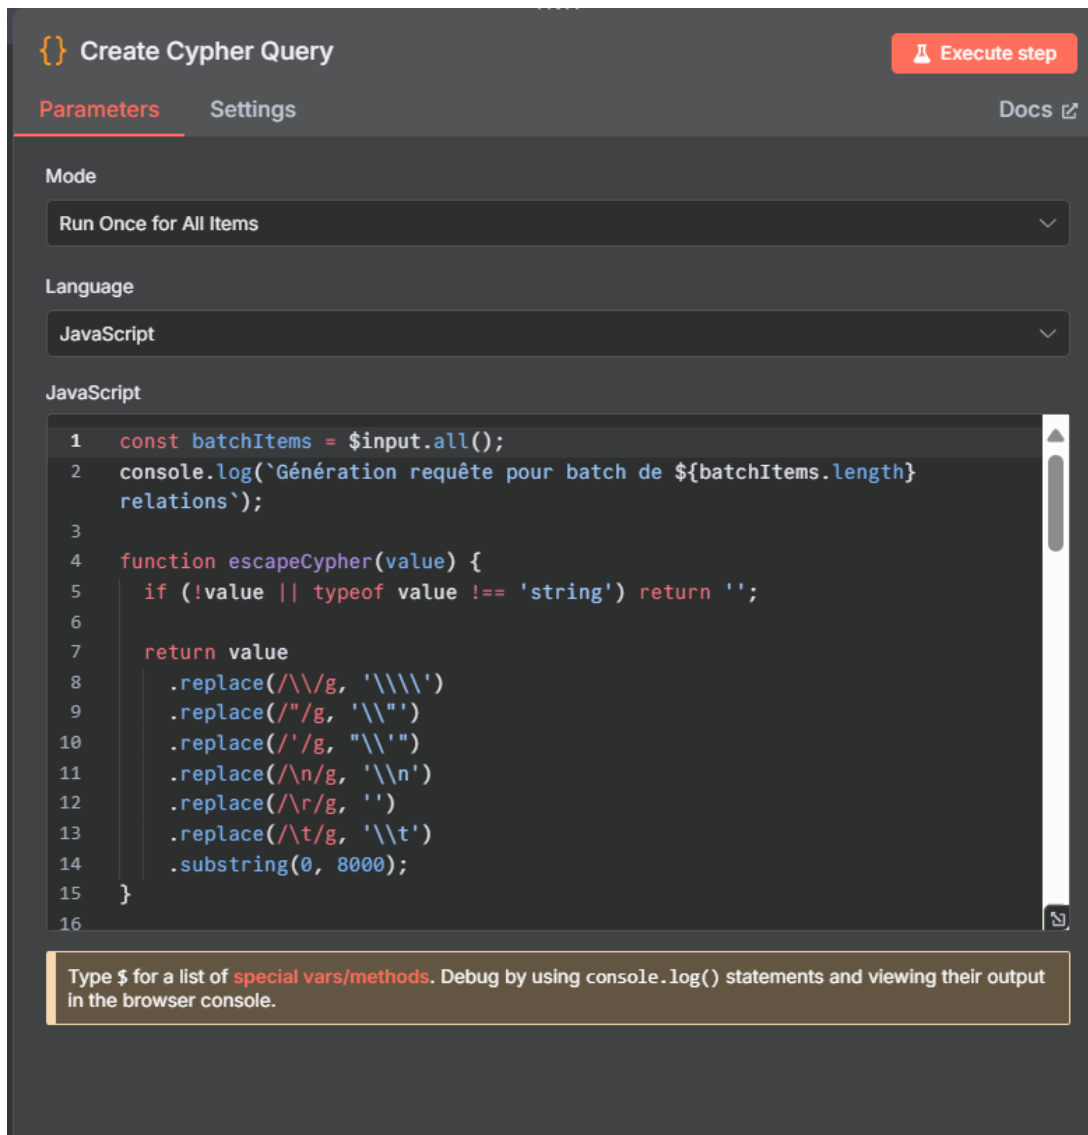


Figure 32: Code Node for creating Relationships query in neo4j

- Creates relationship with dynamic type and properties
- iv. **Consolidated Query Construction:** Combines all MERGE statements into a single executable query:
- Joins all individual MERGE statements with newlines
 - Adds comprehensive RETURN clause for verification
 - Includes batch metadata (size, timestamp, total relations)
- (e) **ExecuteQuery graphDb (Neo4j Bulk Relationship Creation)** Executes the dynamically generated Cypher query from the previous node using `{{ $json.cypherQuery }}` parameter. The executed query structure:
- ```

MERGE (src0:TestGraph1 {Source: "File1.java", section: "algorithms"})
MERGE (tgt0:TestGraph1 {Source: "File2.md", section: "documentation"})
MERGE (src0)-[rel0:'documents']->(tgt0)
ON CREATE SET
 rel0.justification = "File2.md documents the algorithm in File1.java",

```

```

 rel0.source_section = "algorithms",
 rel0.target_section = "documentation",
 rel0.extractedAt = "2024-01-15T10:30:00.000Z"
ON MATCH SET
 rel0.justification = "File2.md documents the algorithm in File1.java",
 rel0.extractedAt = "2024-01-15T10:30:00.000Z"
[... continues for all relationships in batch ...]

RETURN src0.Source as source0, type(rel0) as relType0,
tgt0.Source as target0,src1.Source as source1,
type(rel1) as relType1, tgt1.Source as target1,
 [... continues for all relationships ...]

```

**ExecuteQuery graphDb** Execute step

**Parameters** **Settings** Docs

Credential to connect with  
Neo4j account Cloud

Resource  
Graph Database

Operation  
Execute Query

Index Name  
vector

Cypher Query  
{{ \$json.cypherQuery }}

Figure 33: Execute Query node configuration for creating relationships

### Summary of inserted data:

- `source` → Source file name
- `target` → Target file name
- `relationshipType` → Relationship type
- `justification` → AI model reasoning
- `source_section`, `target_section` → Corresponding sections
- `extractedAt` → Generation date



## 13 Workflow No. 2: Graph RAG Agent

### 13.1 Architecture

Chat Message → AI Agent → OpenAI Chat Model → Postgres Memory → Neo4j Graph Database (Tool)

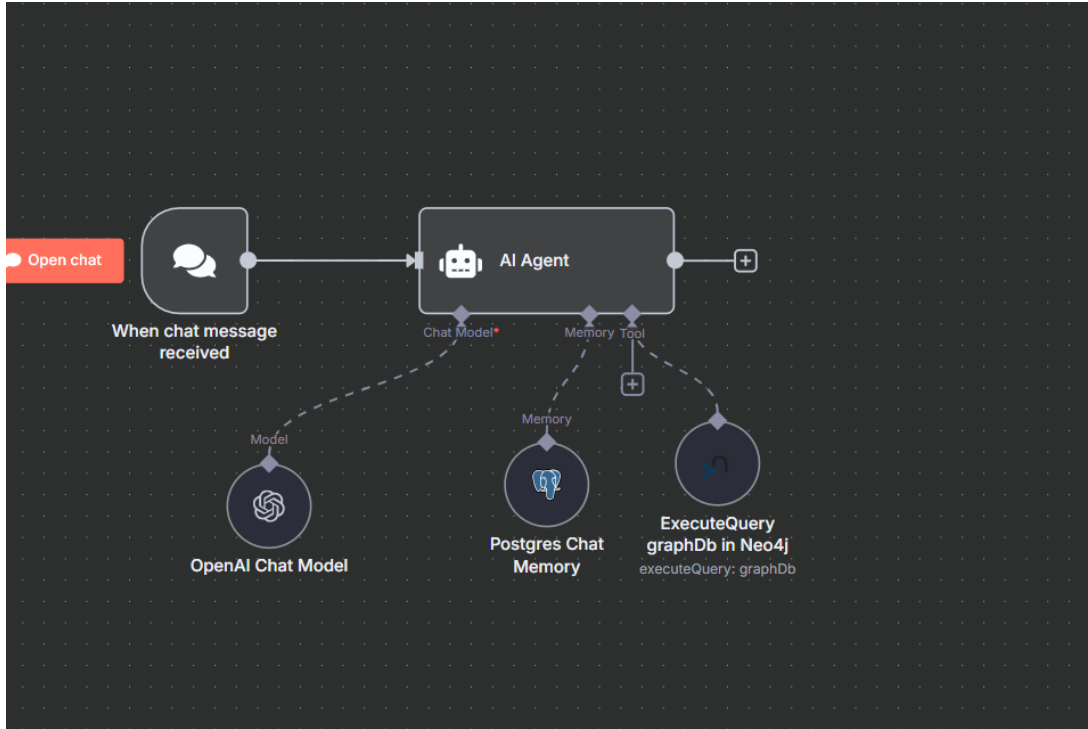


Figure 37: Conversational interface for the Neo4j-based Graph RAG Agent

**General Description:** The intelligent agent (AI Agent) receives a user question, queries the Neo4j graph database to retrieve related Java files and documentation, and generates an enriched response explaining the relationships between the files.

### 13.2 Node 2.1: When Chat Message Received

**Role:** Starting point of the workflow. **Function:** Triggers the flow each time a user sends a message in the chat.

### 13.3 Node 2.2: AI Agent

**Role:** Main Graph RAG agent. **Function:** Analyzes the user query, interacts with Neo4j, and composes a response.

Uses:

- A language model to understand and generate the response.
- A memory (Postgres) to store the conversation context.

- A Neo4j tool to retrieve files and their relationships.

### Combination:

- **Chat Model** → understanding and generation
- **Memory** → conversational context
- **Tool (Neo4j)** → retrieve relevant nodes and relationships

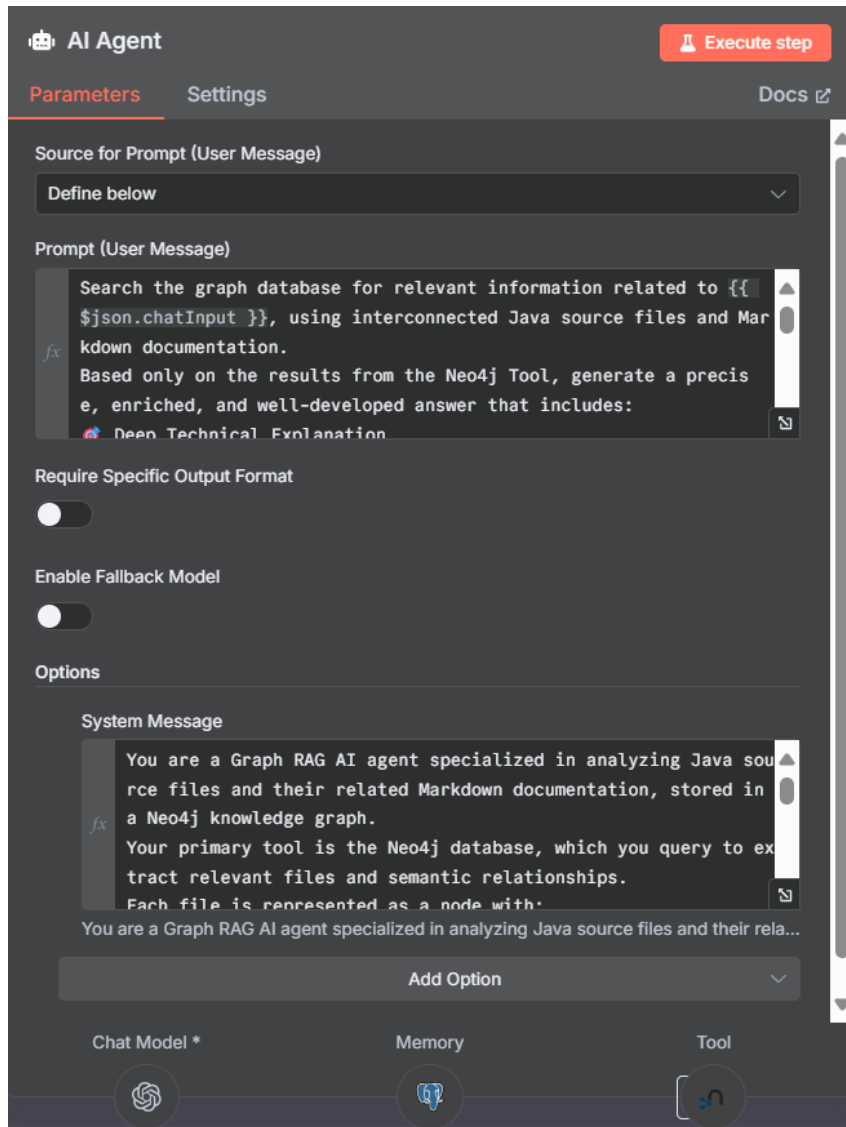


Figure 38: AI Agent configuration connected to Neo4j, Postgres, and OpenAI

### System Message:

You are a Graph RAG AI agent specialized in analyzing Java source files and their related Markdown documentation, stored in a Neo4j knowledge graph. Your primary tool is the Neo4j database, which you query to extract relevant files and semantic relationships.

### Rules:

- Only use data returned from Neo4j (no hallucinations).

- Prioritize relationships: **documents**, **complementary**.
- Include metadata: file name, section, path, extractedAt, snippet of content.
- Respond in the language of the query.

**Expected format:**

- Concise and reasoned response.
- List of files used with their relationship and justification.
- Multi-document citations if needed.

## 13.4 Node 2.3: OpenAI Chat Model

**Role:** Chat GPT 4o mini or Others OpenAI Models and Gemini **Function:** Generates the response based on graph data. Connected to:

- AI Agent via Chat Model.

**Configuration:**

- Add the option *Sampling Temperature* set to 0.0 to eliminate hallucination risk.
- Use the previously created OpenAI credential.

## 13.5 Node 2.4: Postgres Chat Memory

**Role:** Conversational memory. **Function:** Stores conversation history to maintain context.

**Configuration:**

- **Memory Type:** Postgres
- **Connection:** Requires creating a PostgreSQL database credential
- **Required Parameters** from the `.env` file in `self-hosted-ai-starter-kit`:
  - **Host:** postgres (service name in `docker-compose.yml`)
  - **Port:** 5432 (default PostgreSQL port)
  - **Database:** \$POSTGRES\_DB
  - **User:** \$POSTGRES\_USER
  - **Password:** \$POSTGRES\_PASSWORD

Figure 39: PostgreSQL connection configuration

**Example .env content:**

```

POSTGRES_USER=root
POSTGRES_PASSWORD=password
POSTGRES_DB=n8n

N8N_ENCRYPTION_KEY=super-secret-key
N8N_USER_MANAGEMENT_JWT_SECRET=even-more-secret
N8N_DEFAULT_BINARY_DATA_MODE=filesystem

For Mac users running OLLAMA locally
See https://github.com/n8n-io/self-hosted-ai-starter-kit?tab=readme-ov-file#for-mac--apple-silicon-users
OLLAMA_HOST=host.docker.internal:11434

```

Figure 40: Default .env file content

**Connection:** This node is connected to the AI Agent via the **Memory** option.

## 13.6 Node 2.5: ExecuteQuery graphDb in Neo4j

**Role:** Query the Neo4j database to extract files and their relationships. **Function:** Executes the Cypher query to retrieve nodes and relations.

**Cypher Query:**

```

MATCH (n:YourGraphDatabaseName) OPTIONAL MATCH (n)-[r]->(m)
OPTIONAL MATCH (m)-[r2]->(n) RETURN n, r, m, r2

```

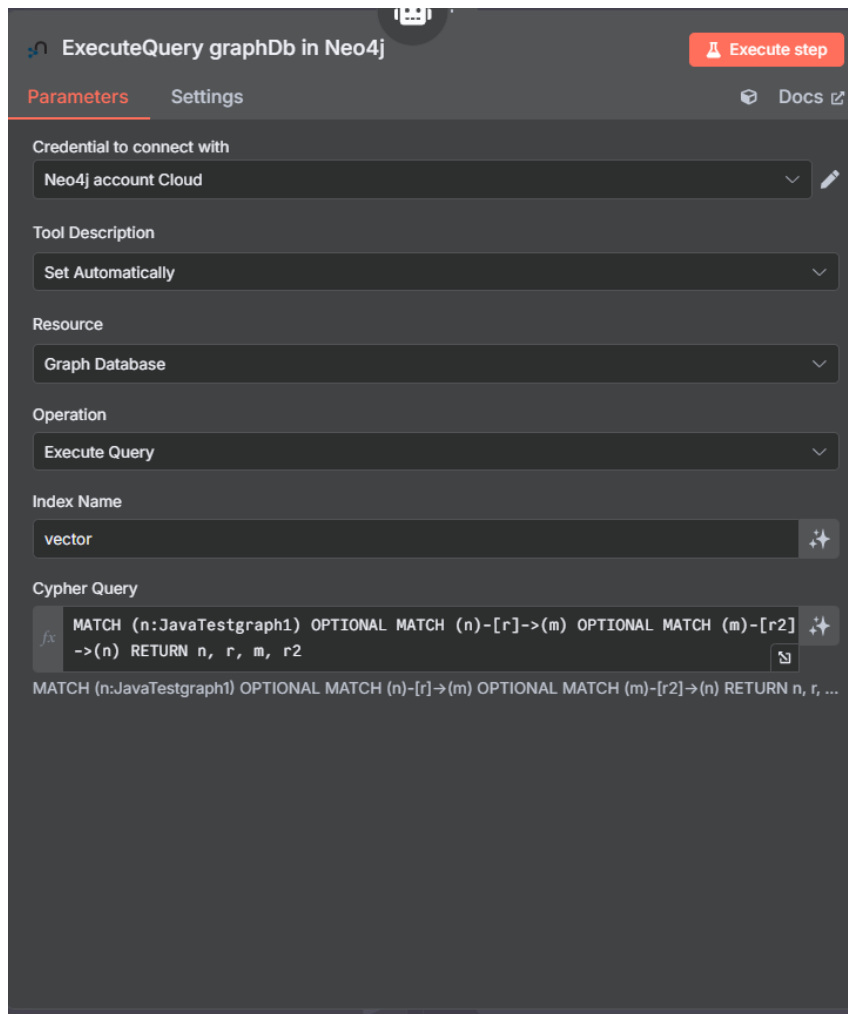


Figure 41: Neo4j node configuration to execute queries

## 14 Test the Graph RAG

After completing the configuration, you can test the Graph RAG by typing a query in the n8n interface and sending it to receive a response.

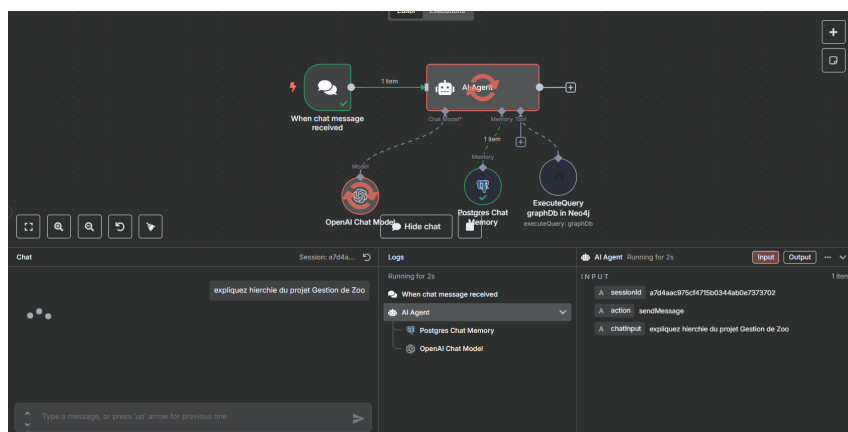


Figure 42: Example query sent in n8n to test the Graph RAG



## 15 Conclusion

This guide helped you set up a complete Graph RAG (Retrieval-Augmented Generation) system using **n8n** for orchestration, **Neo4j** for graph data management, and advanced AI models for semantic analysis and response generation.

### Benefits of the Graph RAG approach:

- **Advanced contextualization:** Semantic relationships between documents allow for deeper understanding and interconnections.
- **Enriched responses:** Responses leverage relationships across multiple documents for coherent answers.
- **Complete traceability:** Each response includes source metadata (file, section, justification).
- **Scalability:** The modular architecture makes it easy to add new document types and relationships.

### Practical applications:

- **Smart documentation:** Automatic analysis of Java codebases and their associated Markdown documentation.
- **Developer assistant:** Helps understand relationships between software components.
- **Technology watch system:** Analyze links across technical and research documents.
- **Enterprise knowledge base:** Smart organization and querying of internal documentation.

**Future extensibility:** The system can be extended to include other file types (PDF, Word, Excel), other graph databases, or specialized AI models based on your organization's needs.

This solution provides a solid foundation for developing advanced AI applications while maintaining full control over your data and leveraging the power of graph-based reasoning.